

Understanding Go Memory Allocation

André Carvalho

 @andresantostc



Developer @ **globo.com**

tsuru 

andrestc.com



THE **LINUX** PROGRAMMING INTERFACE

A Linux and UNIX* System Programming Handbook

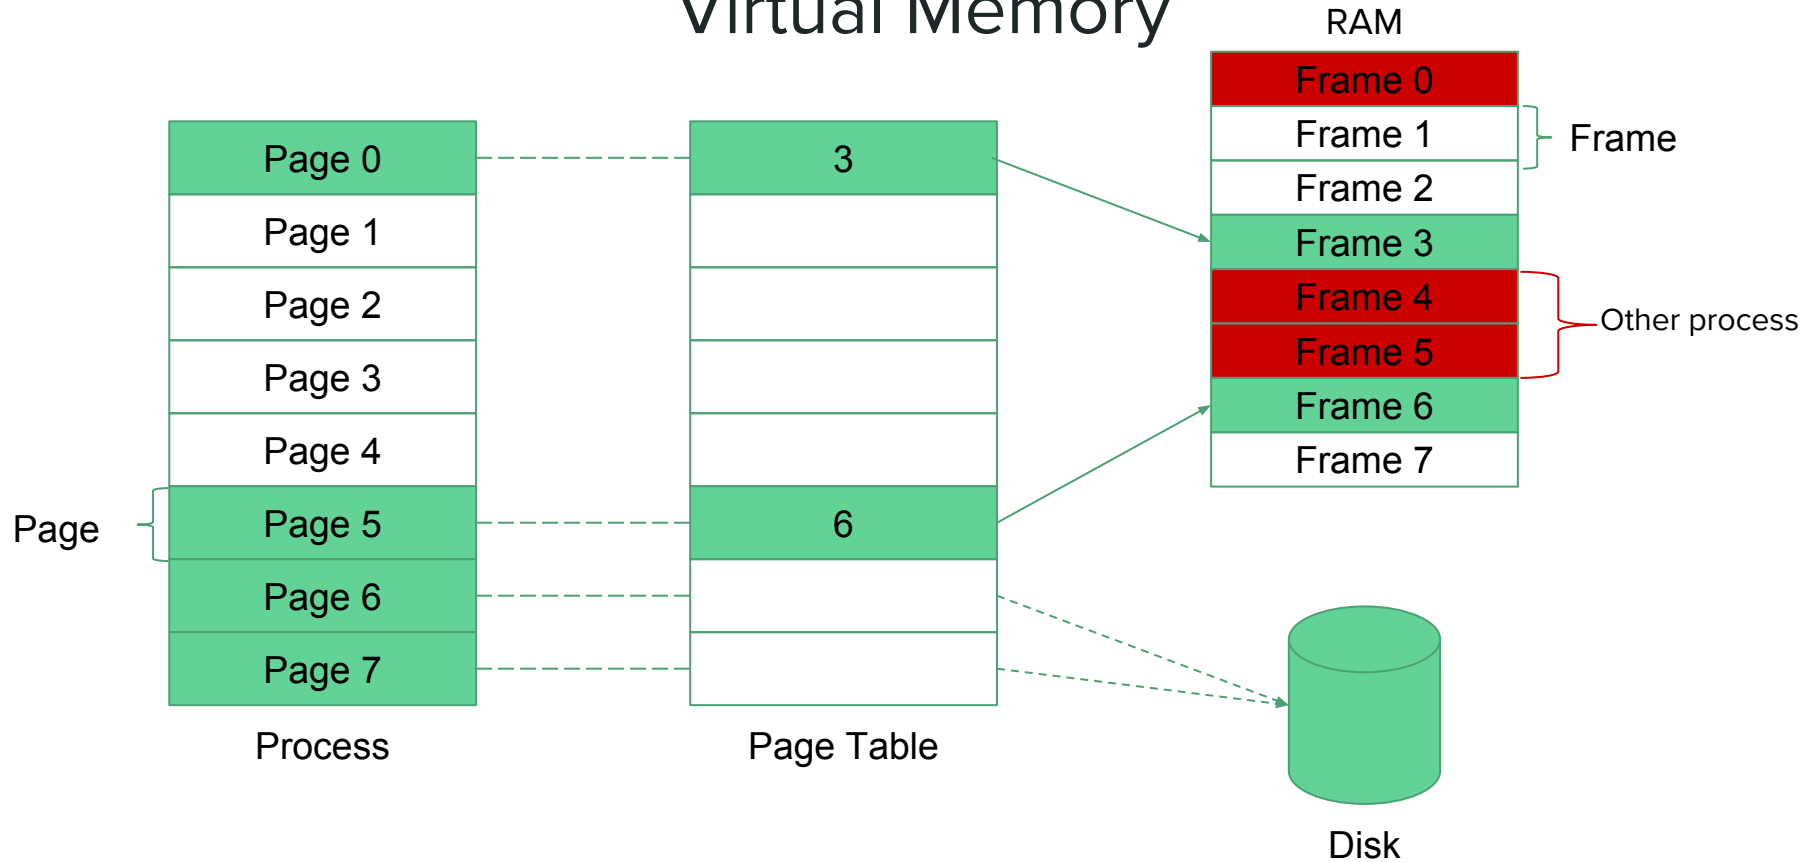
MICHAEL KERRISK



Virtual Memory

- Processes do not read directly from physical memory
 - Security
 - Coordination between multiple processes
- Virtual Memory abstracts that away from the processes
 - Segmentation
 - **Page tables**

Virtual Memory



```
func main() {  
    rand.Seed(time.Now().UnixNano())  
    i := rand.Intn(100)  
    fmt.Printf("%v at %p\n", i, &i)  
    for {  
    }  
}
```

Running two instances at the same time...

\$./vmemory
53 at **0xc420016110**

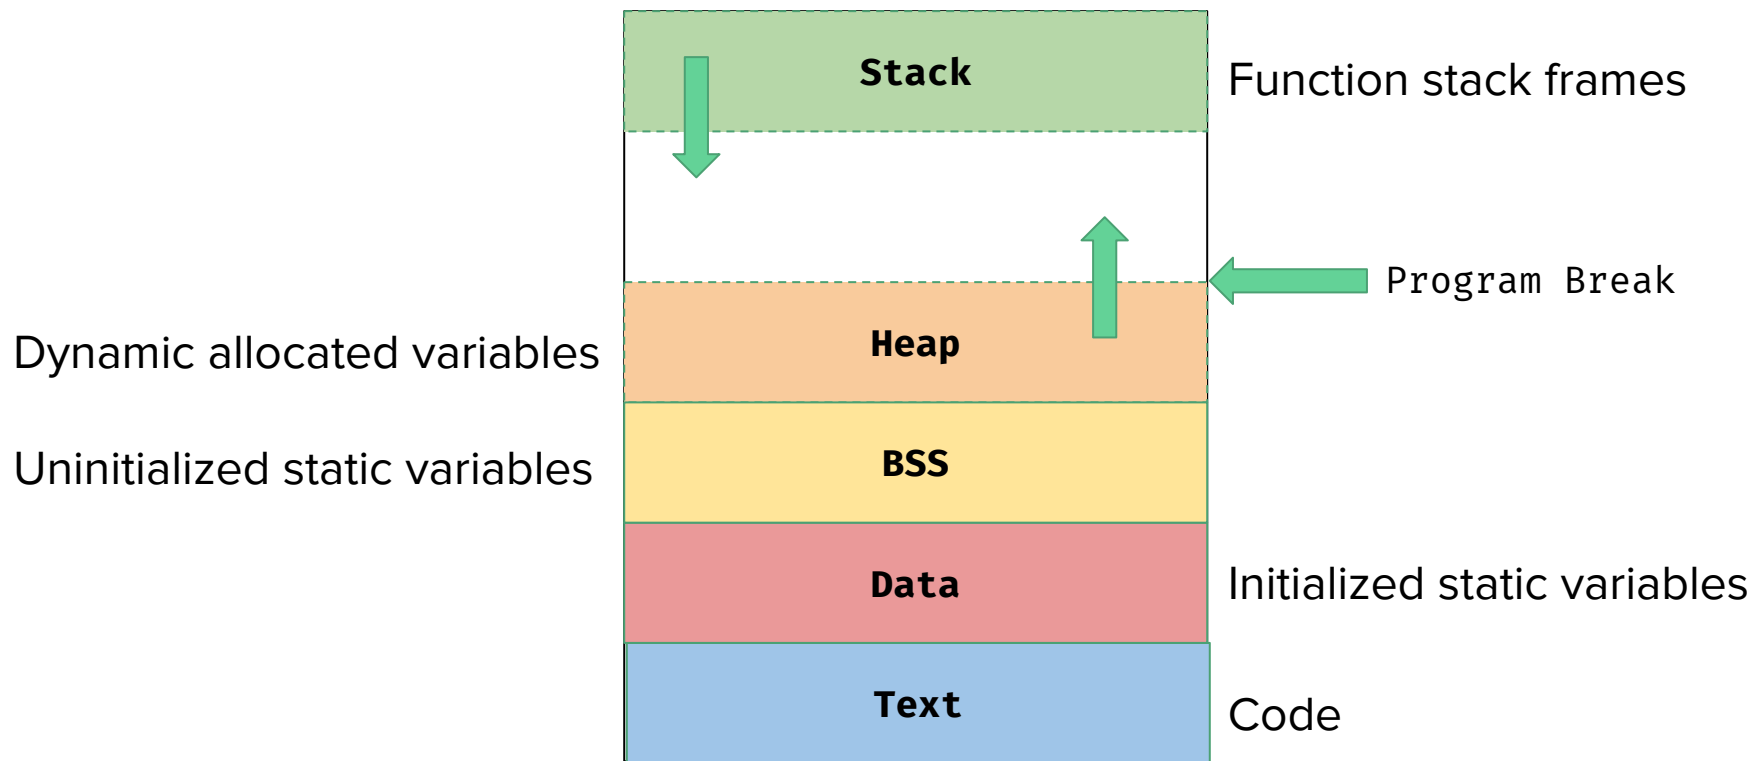
\$./vmemory
68 at **0xc420016110**




Two red arrows originate from a single point at the bottom and point upwards to the underlined virtual address **0xc420016110** in both memory instances above, illustrating that both instances share the same virtual address.

Same virtual address

Process Memory Layout



Stack Allocation



Allocation

```
SP += size;  
return Stack[SP-size];
```

Deallocation

```
SP -= size;
```

Heap Allocation

- For objects with size only known at **runtime**
- C provides **malloc** and **free**
- C++ provides **new** and **delete**
- Go uses **escape analysis** and has **garbage collection**


Minimal Allocator

Minimal Allocator

We need to implement two functions

```
void* malloc(size_t size)
void free(void *ptr)
```


Minimal Allocator



Allocator uses syscalls like `mmap`/`munmap` to allocate/deallocate


Minimal Allocator

Linked list with free objects



Minimal Allocator - Allocating

`malloc(10)`



Minimal Allocator - Allocating

mmap(

0x0000000c00000000, ← **Start Address**

4096, ← **Size**

PROT_WRITE | PROT_READ, ← **Permission**

MAP_PRIVATE | MAP_ANONYMOUS,

...)

↑
Flags




0x000000c000000000

Virtual Address Space

Minimal Allocator - Allocating

`malloc(10)`



Minimal Allocator - Allocating

`malloc(10)`



Minimal Allocator - Allocating



`malloc(10)`




Allocator returns `p`, which points right after the header

Minimal Allocator - Deallocating

`free(p)`



Minimal Allocator - Deallocating



Minimal Allocator

- Can be implemented in a few hundred LOCs
- Issues
 - Fragmentation
 - Corruption
 - Releasing memory back to OS
 - When?
 - How? `munmap`, `madvise`...
 - Multi-thread
 - ...

Go Runtime Allocator

- TCMalloc
- Invoking the Allocator
- Go's Allocator

Thread-Caching Malloc (TCMalloc)

- Originally implemented for the C language by Google
- Served as basis for Go's runtime allocator
- Reduces lock contention for multithreaded programs

TCMalloc

- Each **thread** has a local **cache**
- Two types of allocations
 - **Small** allocations (≤ 32 kB)
 - **Large** allocations
- Manages memory in units called **Spans**
 - Runs of **contiguous** memory pages
 - Metadata is **kept separated from the allocation arena**

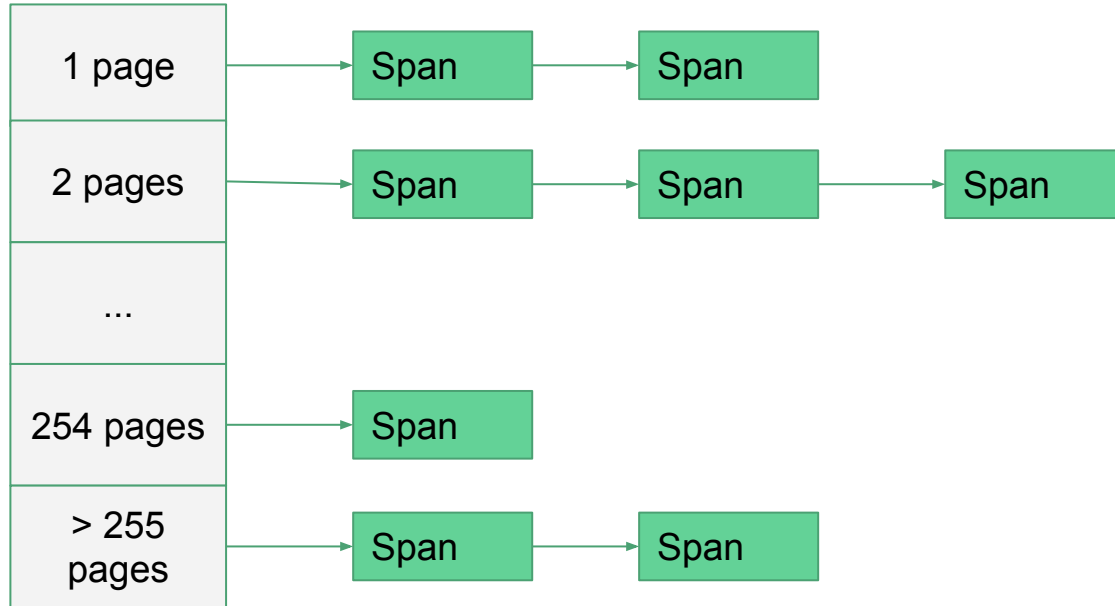
TCMalloc - Large Allocations

- Served by the central heap
- Requested size is rounded up to number of pages (4kB)

`malloc(34 kB) ⇒ malloc(36 kB) ⇒ 9 pages`

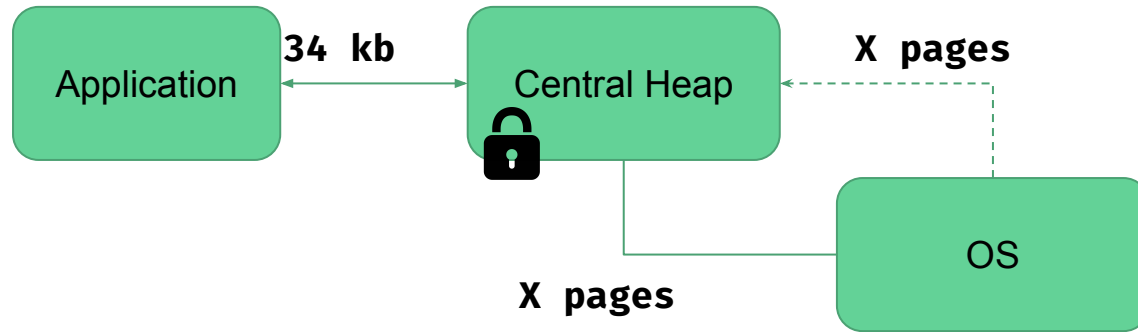
`malloc(33 kB) ⇒ malloc(36 kB) ⇒ 9 pages`

TCMalloc - Large Allocations



Central Heap

TCMalloc - Large Allocations




TCMalloc - Small Allocations

- Served by the local thread cache
- Requested size is rounded up to one of the size classes

`malloc(4 bytes) ⇒ malloc(8 bytes)`


`malloc(6 bytes) ⇒ malloc(8 bytes)`

TCMalloc - Small Allocations




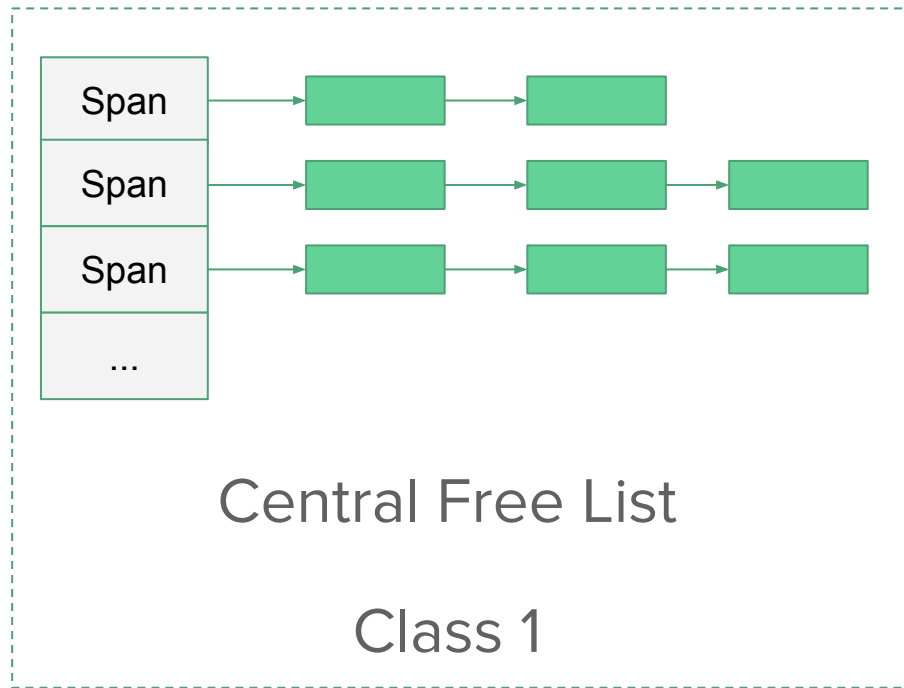

Local Thread Cache

TCMalloc - Small Allocations




Local Thread Cache


TCMalloc - Small Allocations



TCMalloc - Small Allocations



Local Thread Cache



Central Free List

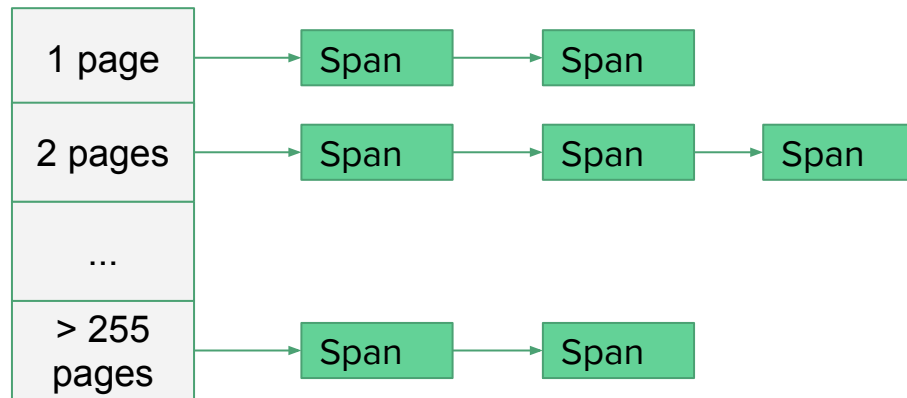
Class 1

TCMalloc - Small Allocations




Central Free List

Class 1




Central Heap

TCMalloc - Small Allocations



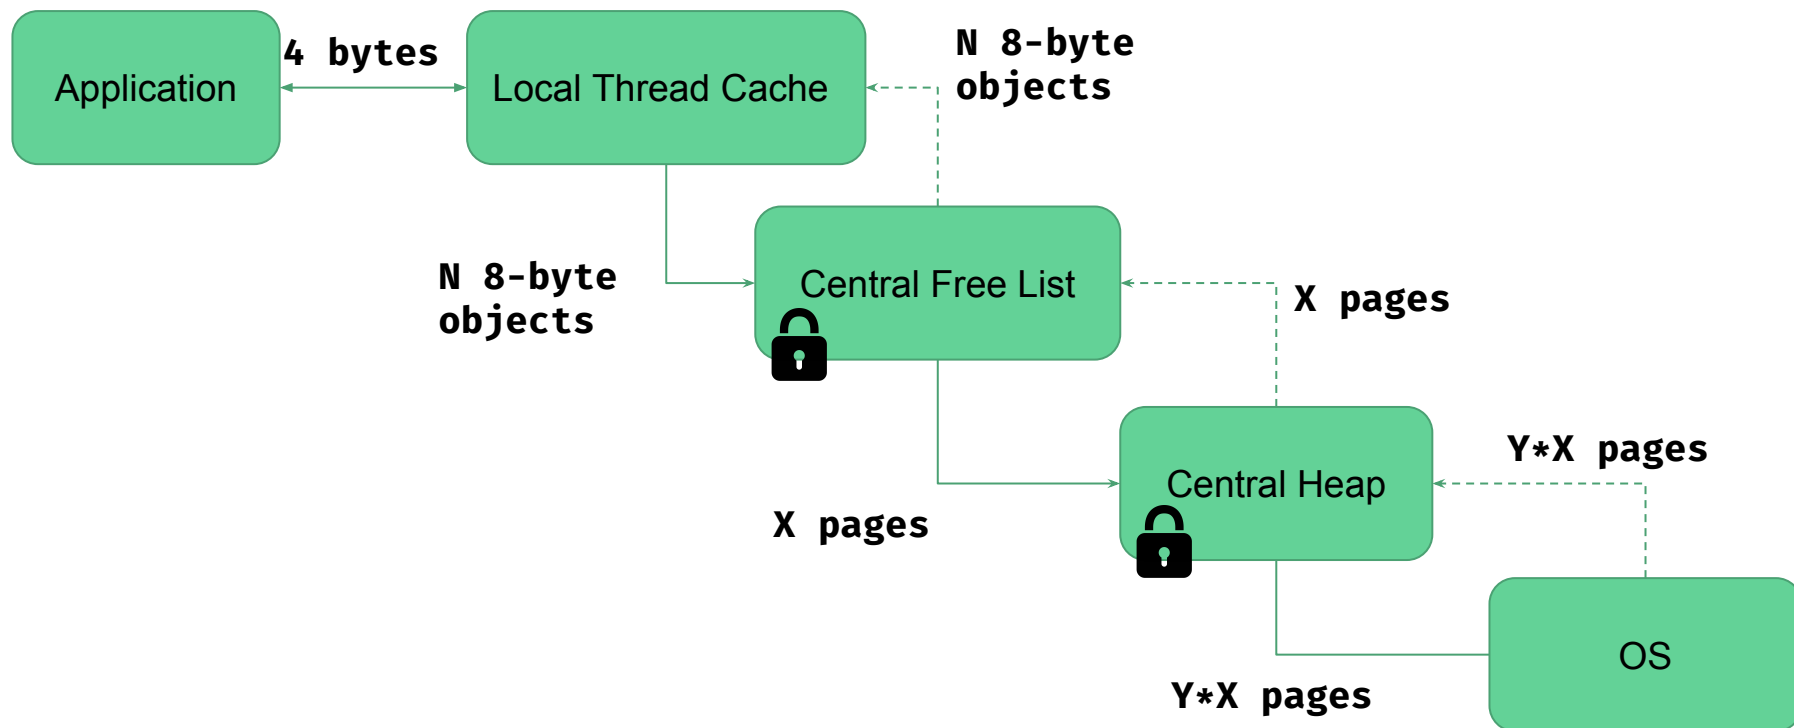
Central Free List

Class 1

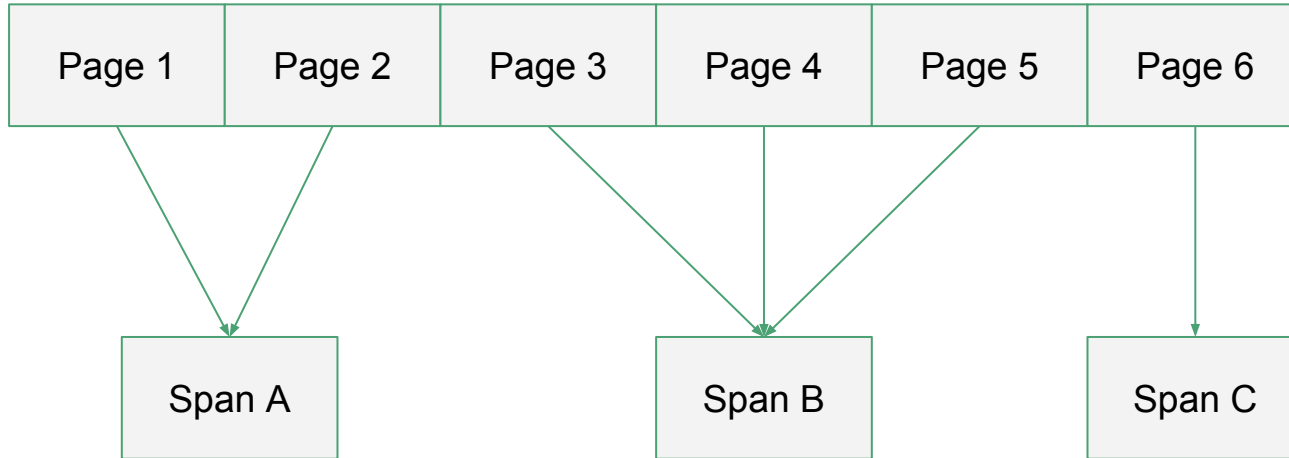


Central Heap

TCMalloc - Small Allocations



TCMalloc - Deallocation




TCMalloc - Deallocation

free()




TCMalloc - Deallocation

free()




Small object



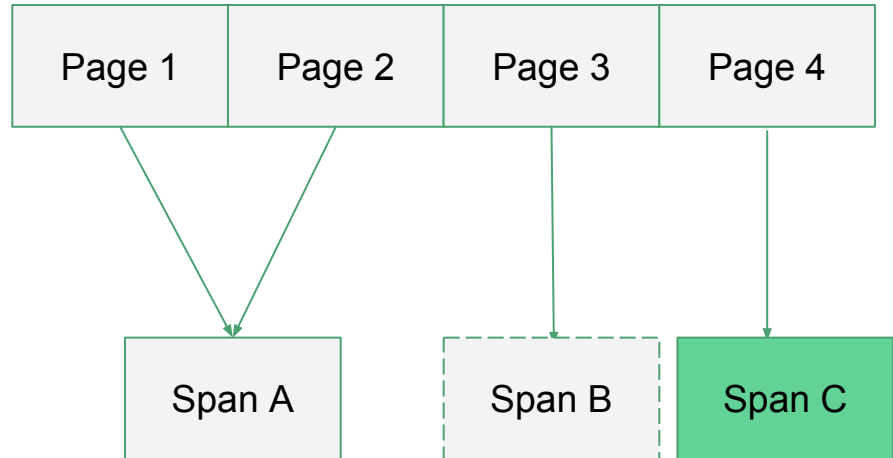
Local Thread Cache

TCMalloc - Deallocation

free()




Large object

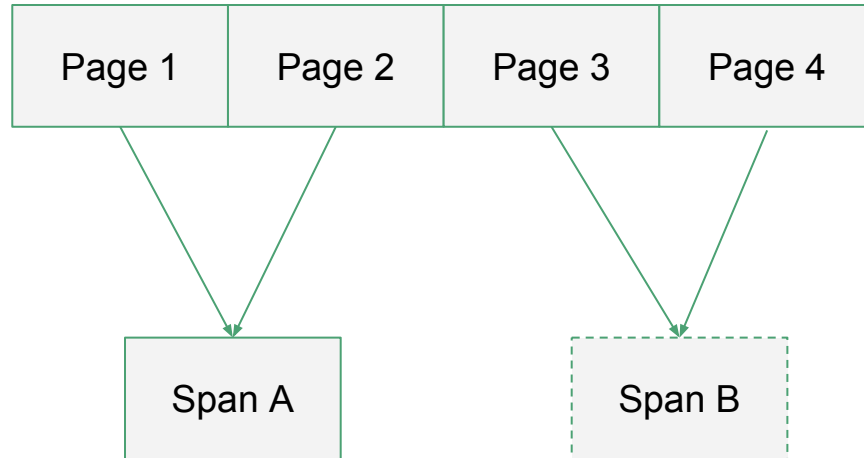


TCMalloc - Deallocation

free()




Large object

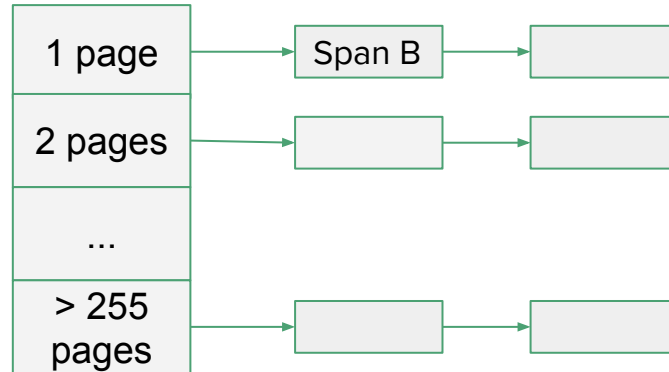


TCMalloc - Deallocation

free()




Large object




Central Heap

TCMalloc - Deallocation

free()



Large object



Central Heap

Go Runtime Allocator

- TCMalloc
- Invoking the Allocator
- Go's Allocator

```
package main
```

```
func main() {  
    f()  
}
```

```
//go:noinline  
func f() *int {  
    i := 10  
    return &i  
}
```

```
package main
```

```
func main() {
```

```
    f()
```

```
}
```

```
//go:noinline
```

```
func f() *int {
```

```
    i := 10
```

```
    return &i
```

```
}
```

```
$ go build -gcflags "-m -m" main.go
```

```
# command-line-arguments
```

```
./main.go:8:6: cannot inline f: marked
```

```
go:noinline
```

```
./main.go:3:6: cannot inline main: non-leaf
```

```
function
```

```
./main.go:10:9: &i escapes to heap ←
```

```
./main.go:10:9:    from ~r0 (return) at
```

```
./main.go:10:2
```

```
./main.go:9:2: moved to heap: i ←
```

```
$ go tool compile -S main.go
```

```
...
```

```
0x001d 00029 (main.go:9)
```

```
LEAQ  type.int(SB), AX
```

```
0x0024 00036 (main.go:9)
```

```
MOVQ  AX, (SP)
```

```
0x0028 00040 (main.go:9)
```

```
PCDATA      $0, $0
```

```
0x0028 00040 (main.go:9)
```

```
CALL  runtime.newobject(SB)
```

```
...
```

```
$ go tool compile -S main.go
```

```
...
```

```
0x001d 00029 (main.go:9)      LEAQ  type.int(SB), AX
```

```
0x0024 00036 (main.go:9)      MOVQ  AX, (SP)
```

```
0x0028 00040 (main.go:9)      PCDATA      $0, $0
```

```
0x0028 00040 (main.go:9)      CALL  runtime.newobject(SB)
```

```
...
```

```
func newobject(typ *_type) unsafe.Pointer {  
    return mallocgc(typ.size, typ, true)  
}
```


Go Runtime Allocator

- TCMalloc
- Invoking the Allocator
- Go's Allocator

Go's Allocator

- Based of TCMalloc
- Garbage Collector
 - Tightly coupled with the allocator
 - Makes hard (impossible?) to replace with other implementations
- Three types of allocations
 - Tiny Allocations (size < 16 bytes, no pointers)
 - Small Allocations (size <= 32 kbytes)
 - Large Allocations

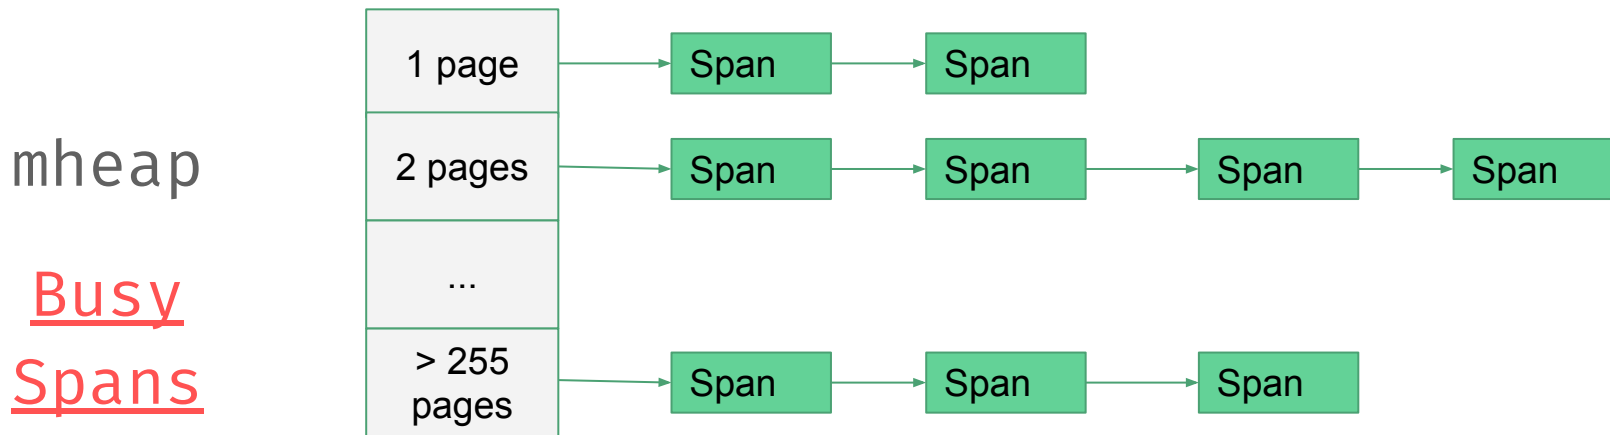
Go's Allocator - Sweeping

Garbage Collector \Rightarrow Concurrent mark and sweep



1. Scan all objects
2. **Mark** objects that are live
3. **Sweep** objects that are not live
 - a. In background
 - b. In response to allocations

Go's Allocator - Large Allocations

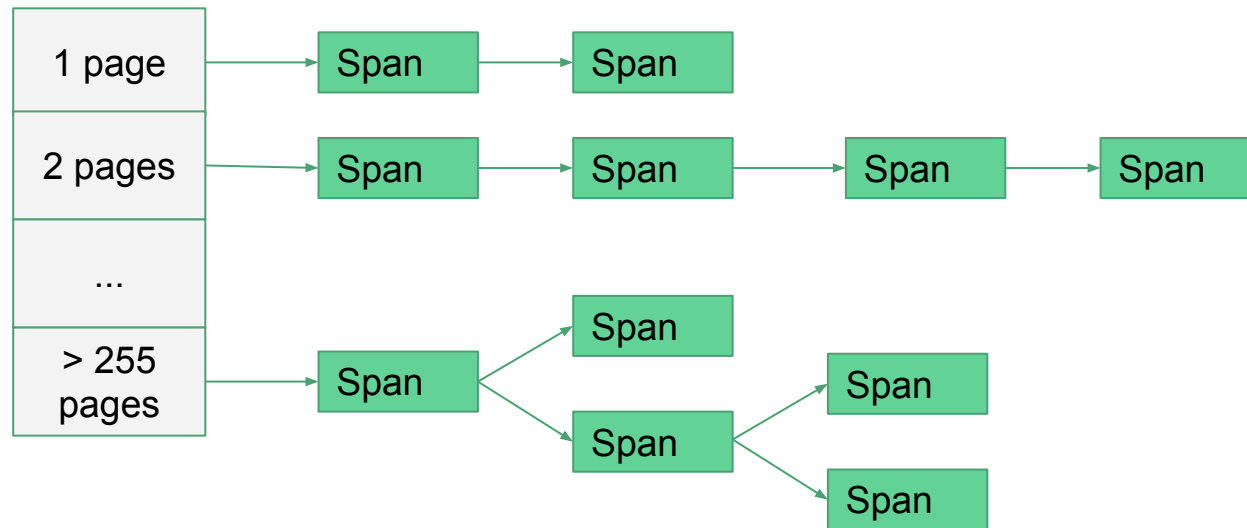


Before allocating, mheap **sweeps** the requested number of pages

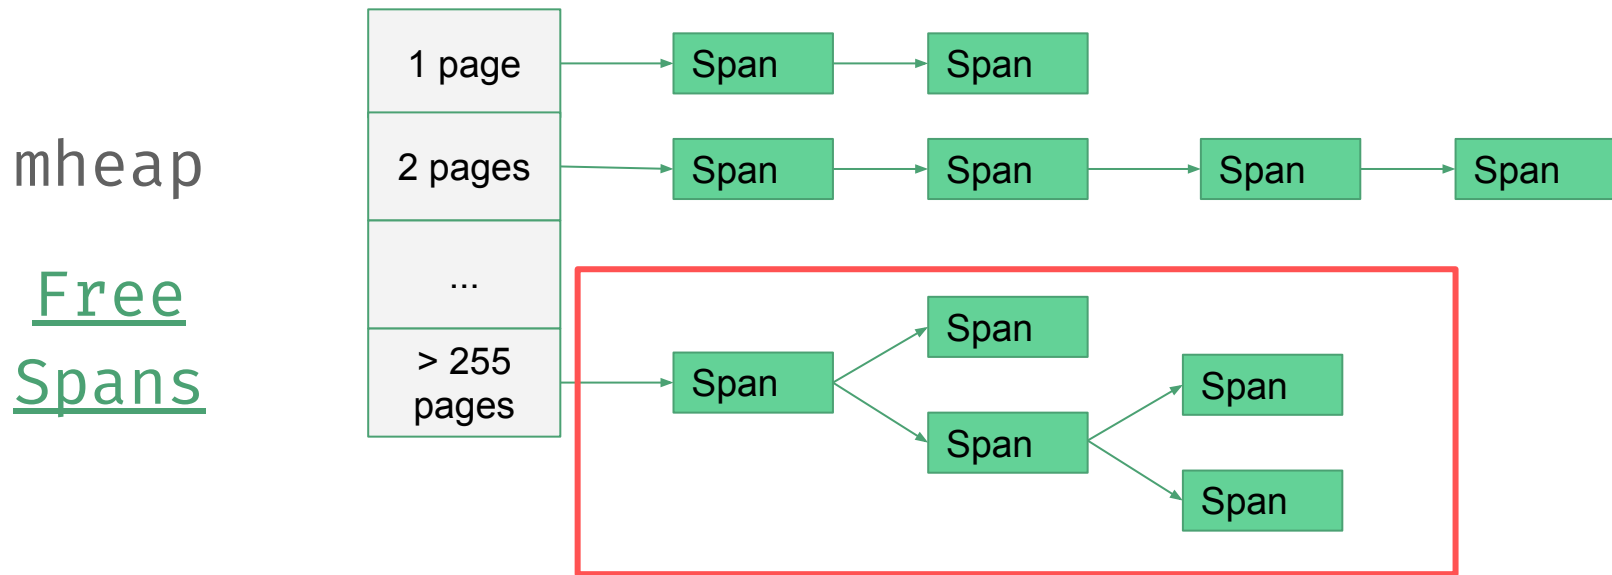
Go's Allocator - Large Allocations

mheap

Free
Spans



Go's Allocator - Large Allocations




mtreap \Rightarrow randomized binary tree

Go's Allocator - Large Allocations

After allocating, depending on the total amount of live memory...


The goroutine may perform additional work for the GC!

Go's Allocator - Small Allocations



Each logical processor (P) has a local cache (mcache)

Go's Allocator - Small Allocations

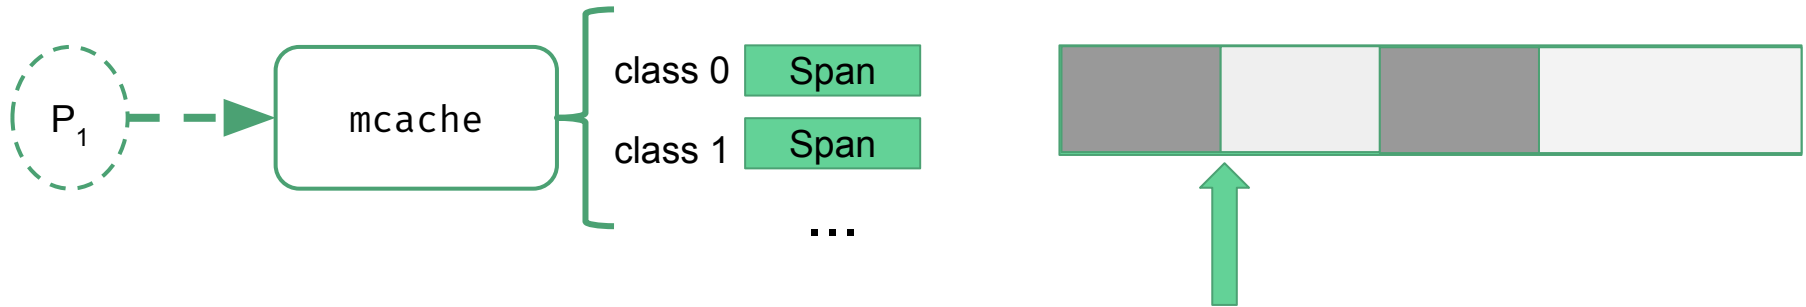


Each mcache maintains a span for each **size class**

Go's Allocator - Small Allocations

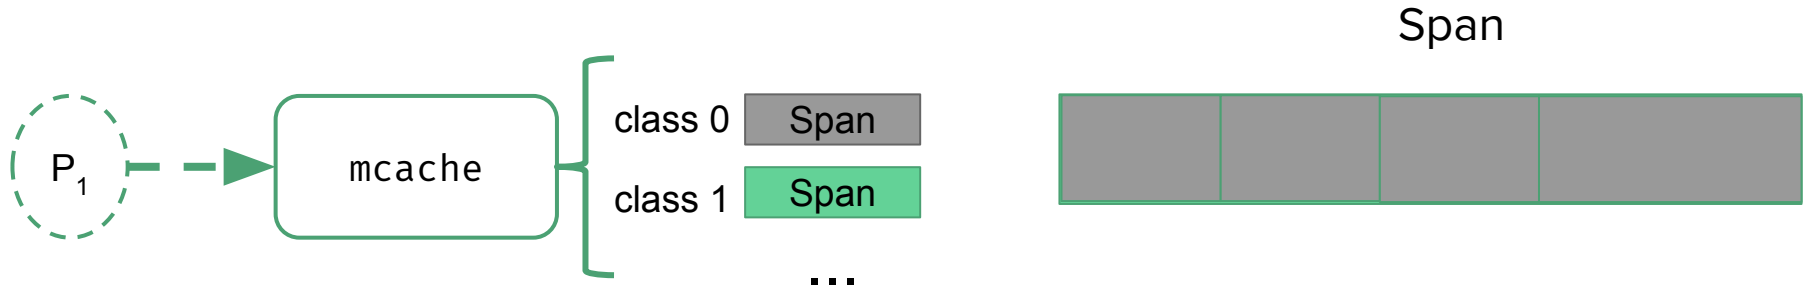
class	bytes/obj	bytes/span	objects
1	8	8192	1024
2	16	8192	512
3	32	8192	256
4	64	8192	170
...			
65	28672	57344	2
66	32768	32768	1

Go's Allocator - Small Allocations



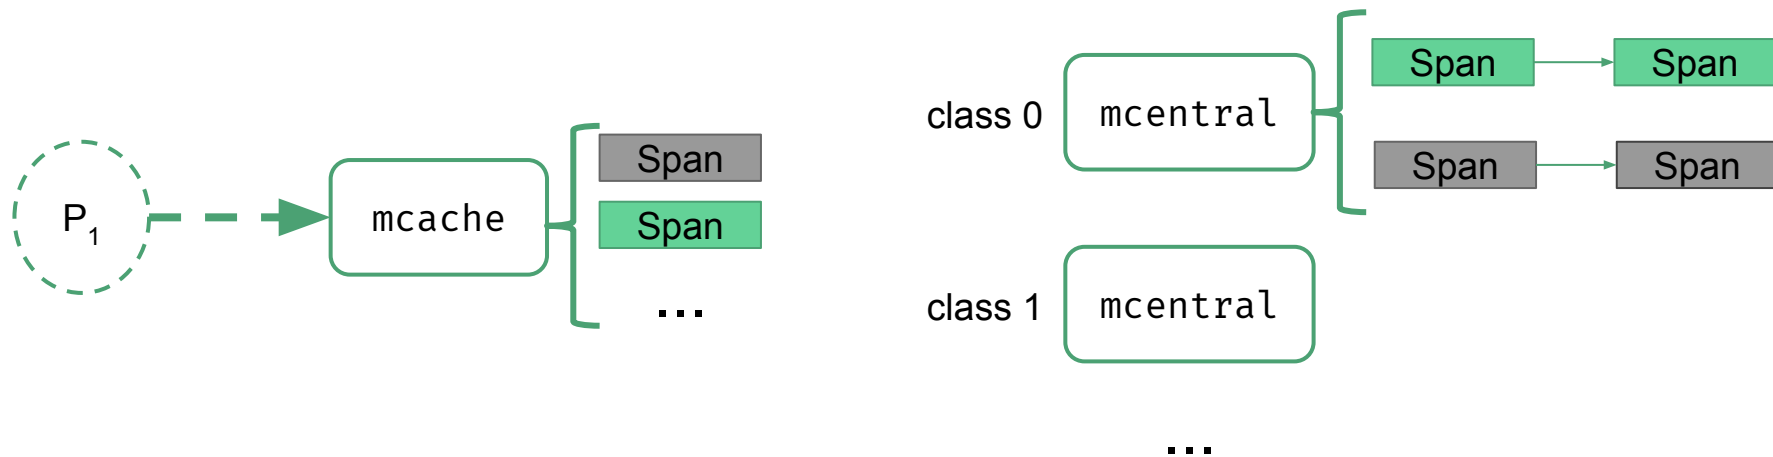
mcache returns the address for a free object on the span

Go's Allocator - Small Allocations



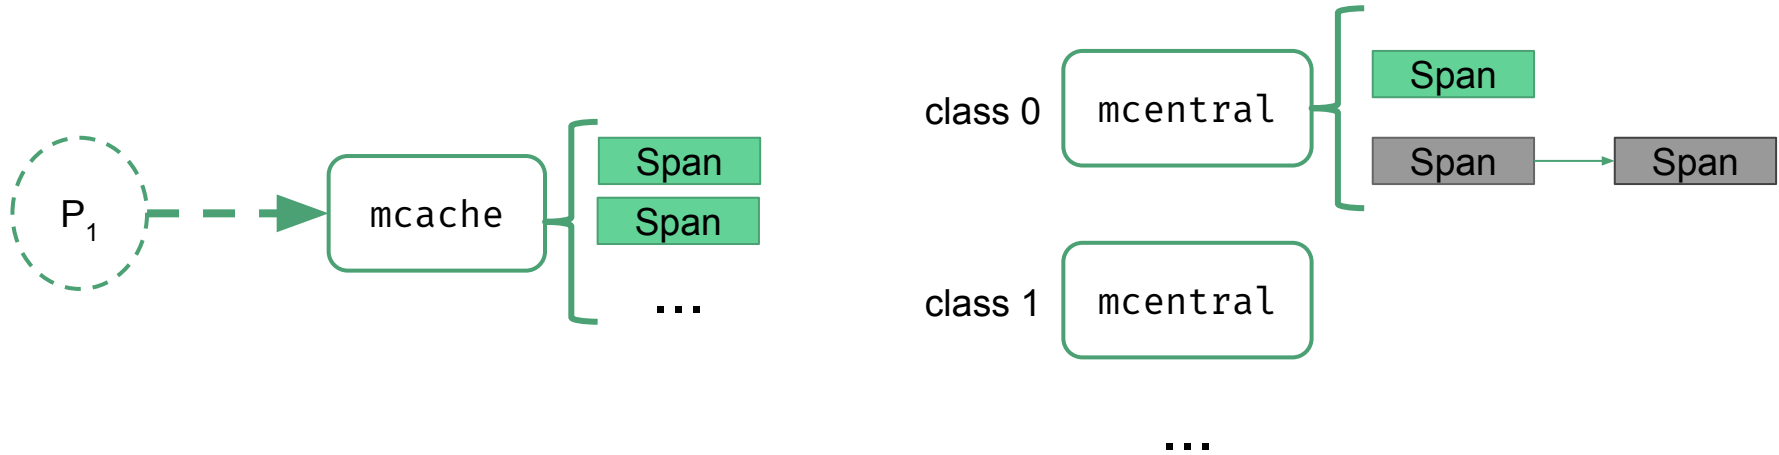
`mcache` request a new span from `mcentral` for this **size class**

Go's Allocator - Small Allocations



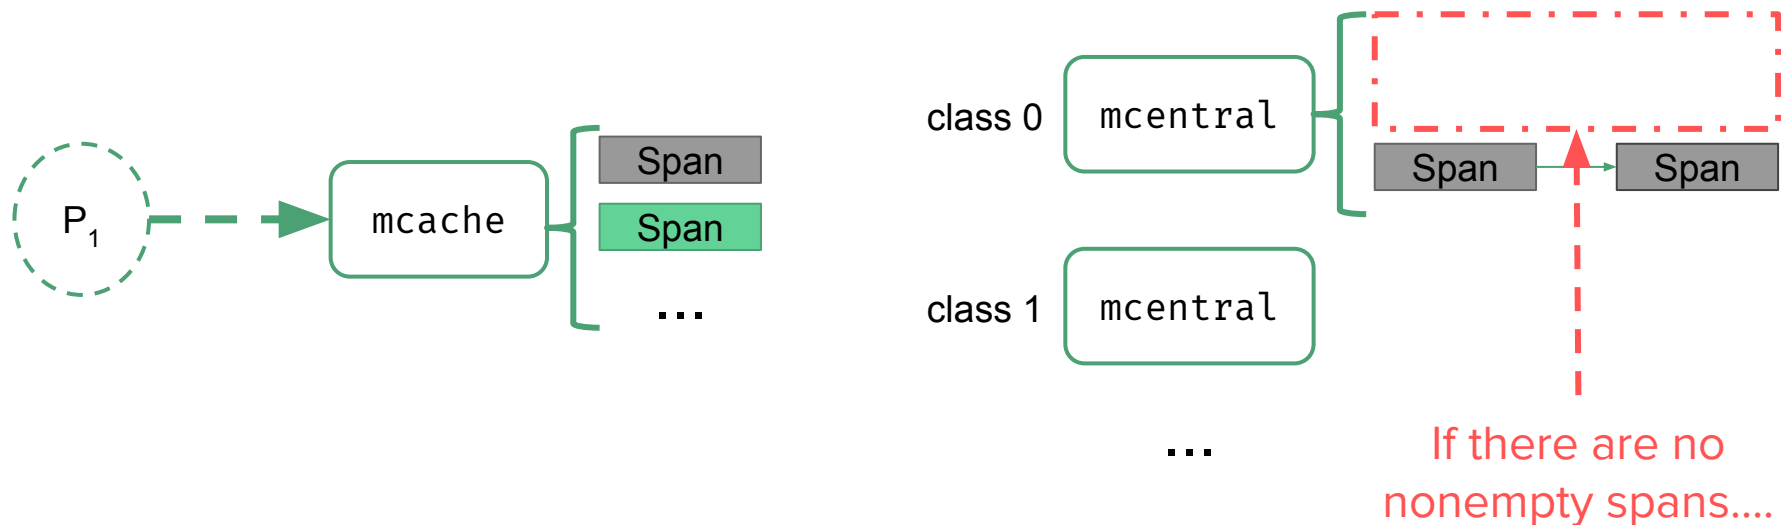
Each `mcentral` has two linked lists, empty and nonempty spans

Go's Allocator - Small Allocations




Span with free objects will be given to the mcache

Go's Allocator - Small Allocations




`mcentral` will try to sweep existing spans

Go's Allocator - Small Allocations



As a last resort, `mcentral` will ask for a new span from `mheap`


Go's Allocator - Small Allocations



`mcentral` will give this span to `mcache`


Go's Allocator - Tiny Allocations

Allocations for objects with **no** pointers and size < **16 bytes**




The main targets of tiny allocator are **small strings** and standalone escaping variables. On a json benchmark the allocator reduces number of allocations by ~12% and reduces **heap size by ~20%**.

Go's Allocator - Tiny Allocations



- Each **P** keeps a 64-bytes **object** allocated from a span
- Each tiny allocation appends a **subobject**

Go's Allocator - Tiny Allocations



- Grab a new object from the `mcache` \approx small allocation
- Eventually, GC will deallocate the old object

Go's Allocator - Releasing memory to the OS

- Runtime periodically releases memory to the OS
- Releases spans that were swept more than **5 minutes ago**
- In Linux, uses the `madvise(2)` syscall

```
madvise(addr, size, _MADV_DONTNEED)
```

```
stats := runtime.MemStats{}  
runtime.ReadMemStats(&stats)
```



```
type MemStats struct {  
    ...  
    // Heap memory statistics.  
    HeapAlloc uint64  
    HeapSys   uint64  
    HeapIdle  uint64  
    HeapInuse  uint64  
    HeapReleased uint64  
    HeapObjects uint64  
    ...  
}
```

References

1. <http://goog-perftools.sourceforge.net/doc/tcmalloc.html>
2. <https://www.ardanlabs.com/blog/2017/05/language-mechanics-on-stacks-and-pointers.html>
3. <https://gabrieletolomei.wordpress.com/miscellanea/operating-systems/in-memory-layout/>
4. [Lec 10 | MIT 6.172](#) - <https://www.youtube.com/watch?v=p0bc1f6ULxw>
5. <https://faculty.washington.edu/aragon/pubs/rst89.pdf>
6. <http://man7.org/linux/man-pages/man2/mmap.2.html>
7. <http://man7.org/linux/man-pages/man2/madvise.2.html>
8. <https://nostarch.com/tlpi>

Thanks!



andrestc.com



@andresantostc